

The Machine Learning Landscape

Concepts, Tools, and Techniques to Build
Intelligent Systems

Introduction to Machine Learning

- **Definition:**
 - "*Machine learning is the science (and art) of programming computers so they can learn from data.*"
 - "*Field of study that gives computers the ability to learn without being explicitly programmed.*" – Arthur Samuel, 1959
 - "A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E." —Tom Mitchell, 1997
- **Key Components:**
 - Training data
 - Model (e.g., neural networks, random forests)
 - Performance measure (e.g., accuracy, RMSE)

The image shows a close-up of a person's hand pointing their index finger towards a computer monitor. The monitor displays a dark-themed Python script. The script includes several conditional blocks for different mirror operations (MIRROR_X, MIRROR_Y, MIRROR_Z) and logic for selecting specific objects in a 3D scene. The code is written in a clear, readable style with appropriate indentation and comments.

```
mirror_mod = modifier_obj
# mirror object to mirror
mirror_mod.mirror_object

if operation == "MIRROR_X":
    mirror_mod.use_x = True
    mirror_mod.use_y = False
    mirror_mod.use_z = False
elif operation == "MIRROR_Y":
    mirror_mod.use_x = False
    mirror_mod.use_y = True
    mirror_mod.use_z = False
elif operation == "MIRROR_Z":
    mirror_mod.use_x = False
    mirror_mod.use_y = False
    mirror_mod.use_z = True

#selection at the end -add
#mirror_ob.select= 1
#mirror_ob.select=1
context.scene.objects.active = ("Selected" + str(modifier))
mirror_ob.select = 0
# bpy.context.selected_objects[0].select = 0
data.objects[one.name].select = 1
print("please select exactly one object")

-- OPERATOR CLASSES ---

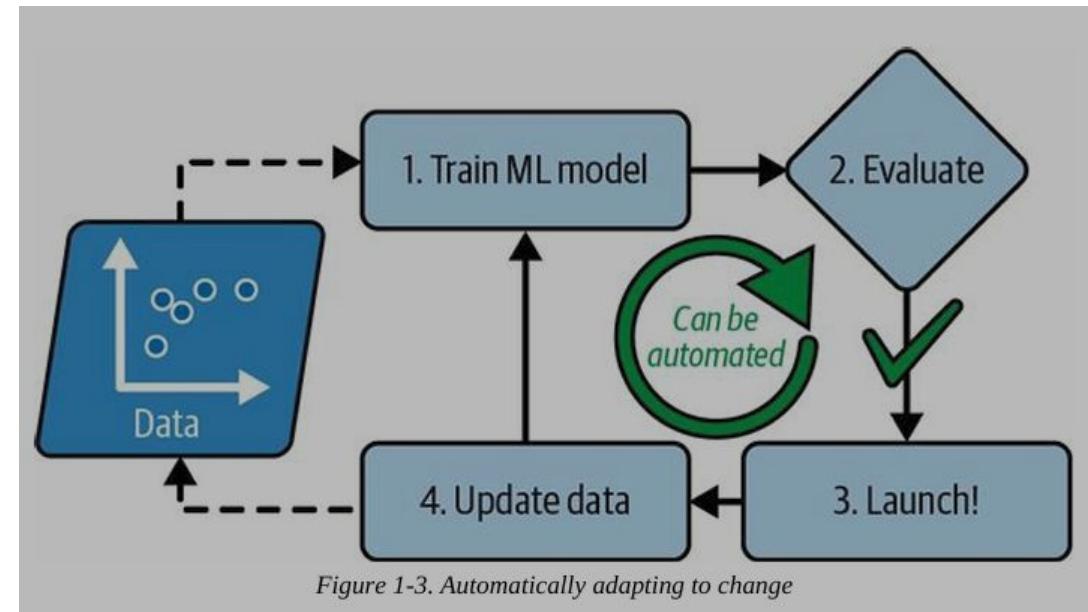
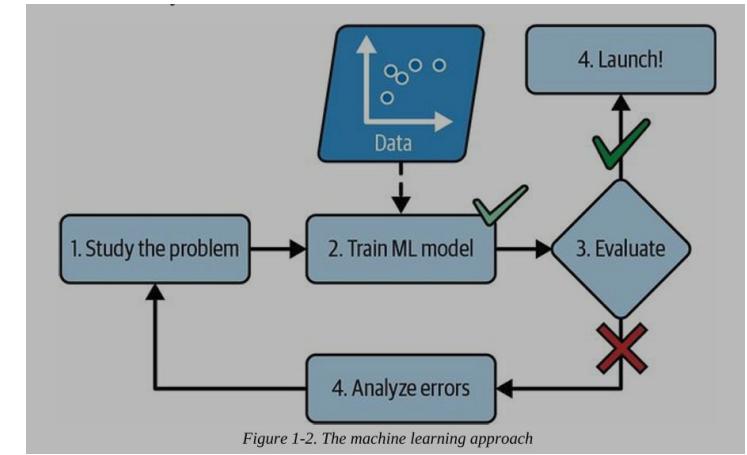
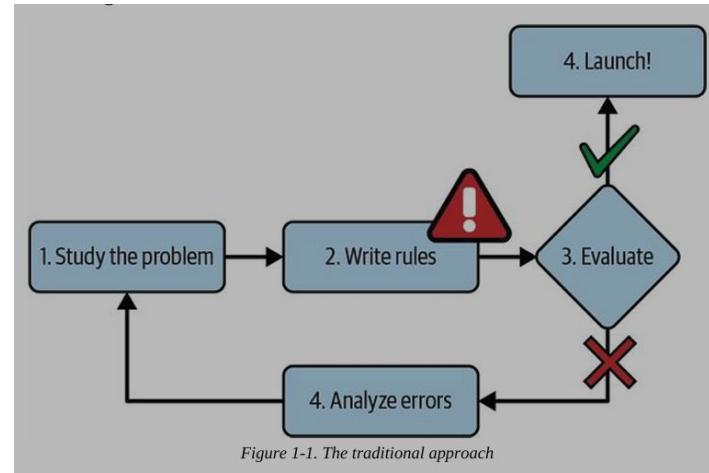
#types.Operator:
# X mirror to the selected
# object.mirror_mirror_x"
# or X"
# context):
# context.active_object is not None
```

Why Use Machine Learning?

- **Advantages:**

- **Automation:** Reduces manual rule-writing (e.g., spam filters).
- **Adaptability:** Learns from new data (e.g., evolving spam tactics).
- **Complex Problems:** Solves tasks with no known algorithm (e.g., speech recognition).
- **Data Mining:** Uncovers hidden patterns (e.g., customer segmentation).

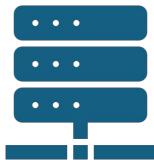
Example: Spam filter vs. manual rules



Examples of ML Applications

- **Image Classification:** CNNs for product sorting (e.g., tumors in scans).
- **Natural Language Processing:** Text classification (e.g., offensive comments, classifying news articles), Text summarization, Creating a chatbot or a personal assistant.
- **Regression:** Forecasting revenue (e.g., linear regression).
- **Anomaly Detection:** Fraud detection (e.g., credit card transactions).
- **Reinforcement Learning:** Game bots (e.g., AlphaGo).
- **Speech Processing:** RNNs, CNNs, or transformers for speech recognition.
- **Clustering:** Segmenting clients based on their purchases
- **Dimensionality Reduction:** Representing a complex dataset in a clear and insightful diagram.
- **Recommender System:** Recommending a product that a client may be interested in, based on past purchases.

Types of ML Systems



Supervised Learning

Labeled data (e.g., spam/ham emails).

Tasks: Classification (spam filter), Regression (house prices).

Unsupervised Learning

Unlabeled data (e.g., customer segmentation).

Tasks: Clustering, dimensionality reduction, anomaly detection, association rule learning.

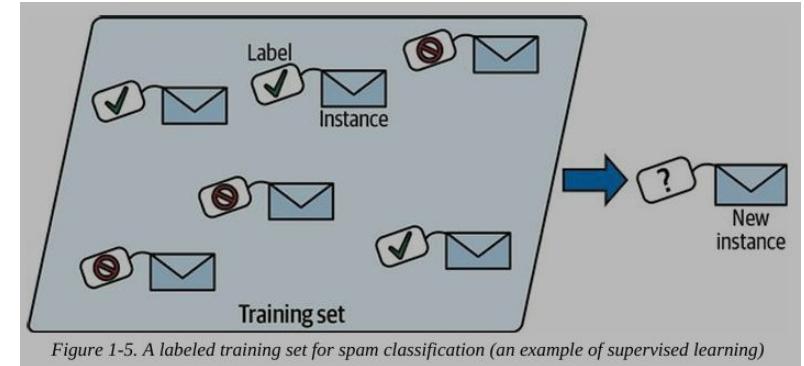


Figure 1-5. A labeled training set for spam classification (an example of supervised learning)

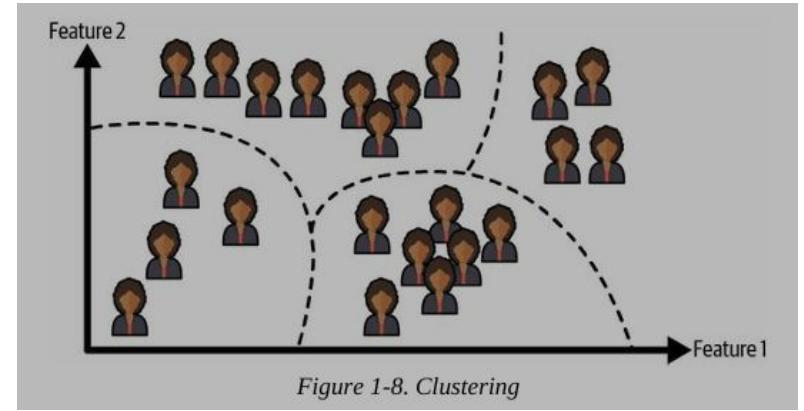


Figure 1-8. Clustering

Types of ML Systems

- **Semi-Supervised Learning**

- Mix of labeled/unlabeled data (e.g., Google Photos face recognition).

- **Self-Supervised Learning**

- Generate fully labeled dataset from a fully unlabeled one (e.g., image inpainting).

- **Reinforcement Learning**

- Agent learns via rewards/penalties (e.g., robot walking).

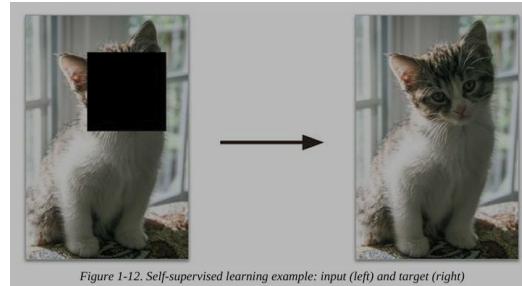
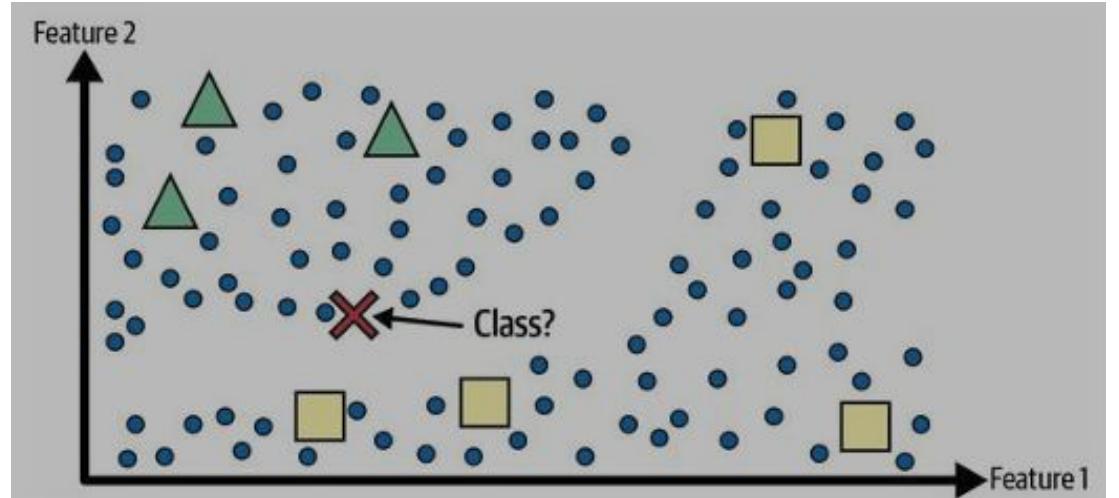
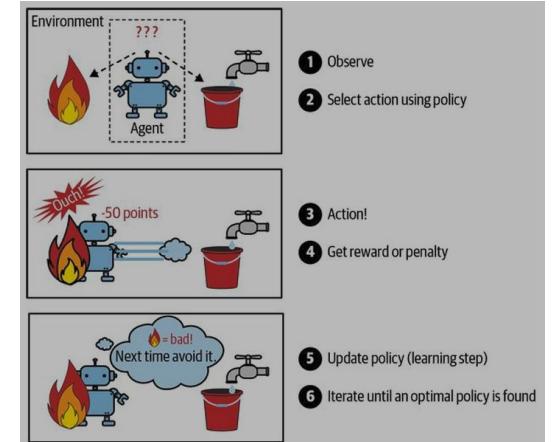


Figure 1-12. Self-supervised learning example: input (left) and target (right)



Batch vs. Online Learning

Batch Learning

- Trained on full dataset (offline).
- **Drawback:** Model rot or data drift (Model's performance tends to decay slowly over time), Require more computing resources.
- **Example:** Classifies pictures of cats and dogs, Making predictions on the financial market

Online Learning

- Incremental updates (system learn about new data on the fly)
- **Drawback:** If bad data is fed to the system, the system's performance will decline.
- **Example:** Stock price predictions.

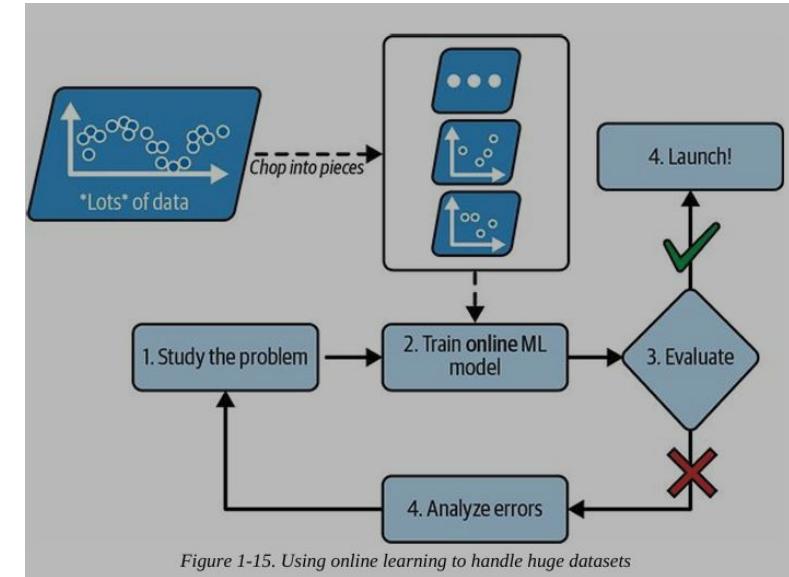


Figure 1-15. Using online learning to handle huge datasets

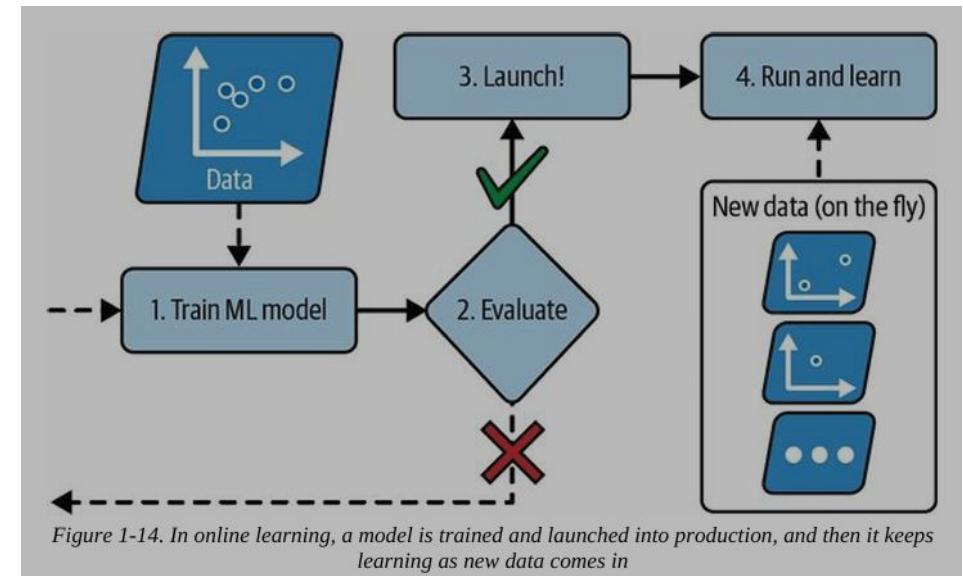
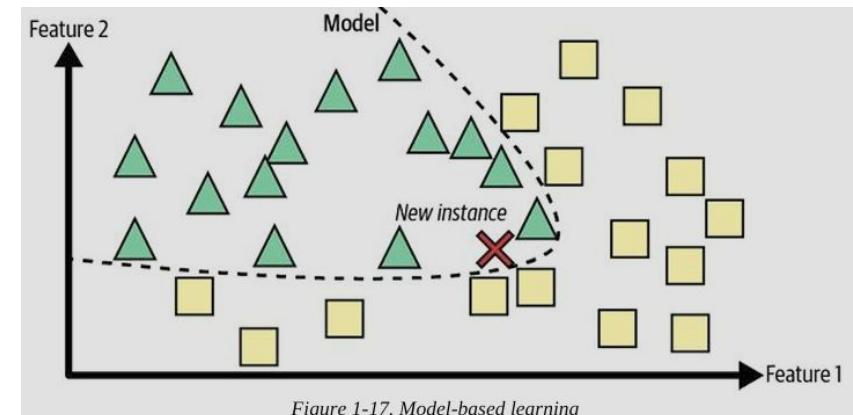
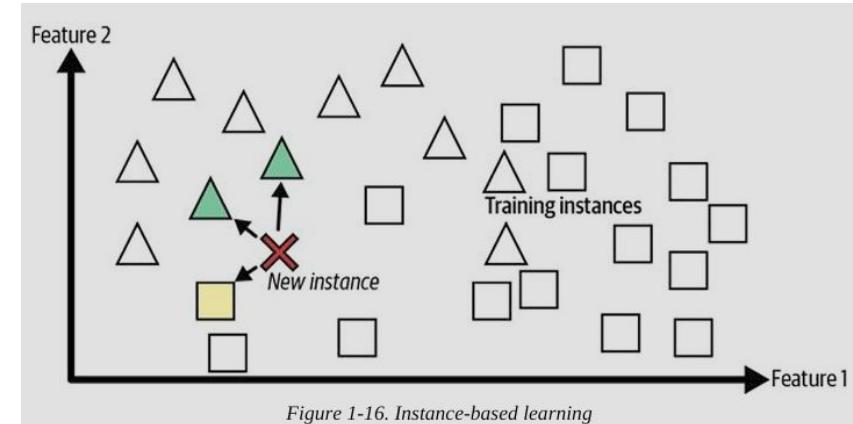


Figure 1-14. In online learning, a model is trained and launched into production, and then it keeps learning as new data comes in

Instance-Based vs. Model-Based Learning

- **Instance-Based:** Memorizes data (e.g., k-nearest neighbors).
 - The system learns the examples by heart, then generalizes to new cases by using a similarity measure to compare them to the learned examples.
- **Model-Based:** Generalizes patterns (e.g., linear regression).
 - Build a model of examples and then use that model to make predictions.
 - **Example:** Predicting life satisfaction:
<https://colab.research.google.com/drive/1vNRge8bHUC3qKg1FA7qjOO2WOFjEqWYg#scrollTo=R37nSnn2pnub>



Main Challenges in ML

- **Insufficient Data** – Small datasets limit performance.
- **Nonrepresentative Data** – Sampling noise (the sample is too small), Sampling bias (very large samples can be nonrepresentative if the sampling method is flawed, e.g., 1936 election poll).
- **Poor-Quality Data** – Errors, outliers, missing values.
- **Irrelevant Features** – Garbage in, garbage out.
- **Overfitting** – Model performs well on the training data, but it does not generalize well.
- **Underfitting** – Model too simple.

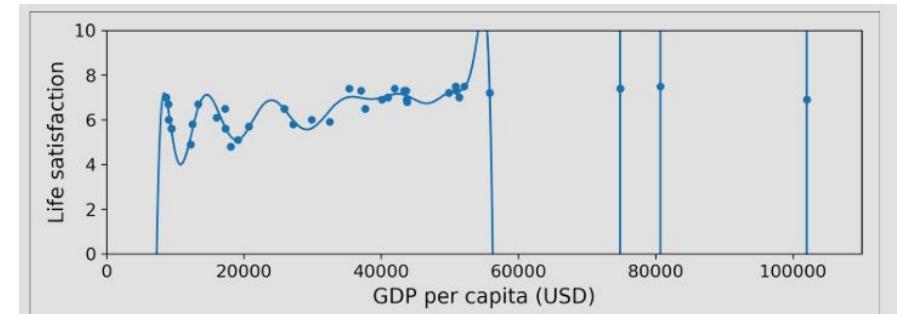
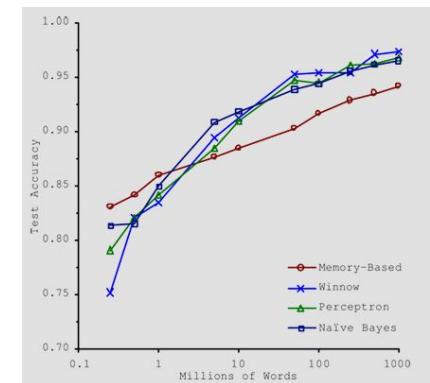
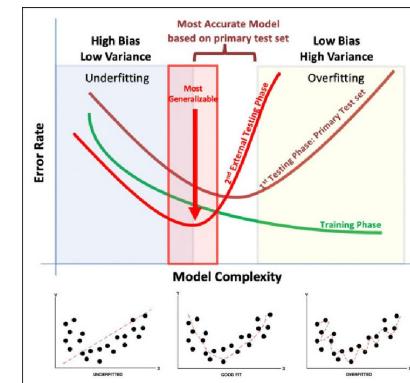


Figure 1-22. Overfitting the training data



Evaluating & Fine-Tuning Models

- **Train/Test Split:** 80% of the data for training and hold out 20% for testing. Compute generalization errors.
- **Cross-Validation:** Reliable performance estimate (e.g., 10-fold CV).
- **Hyperparameter Tuning:**
 - Train multiple models with various hyperparameters on the reduced training set (i.e., the full training set minus the validation set).
 - Select the model that performs best on the validation set.
 - Train the best model on the full training set (including the validation set).
 - Evaluate this final model on the test set to get an estimate of the generalization error.

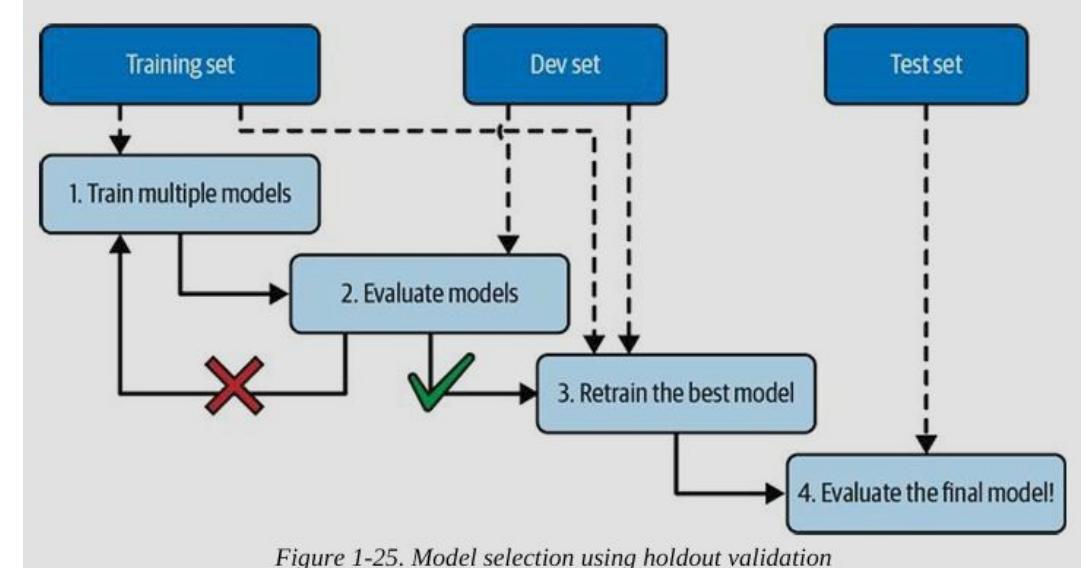
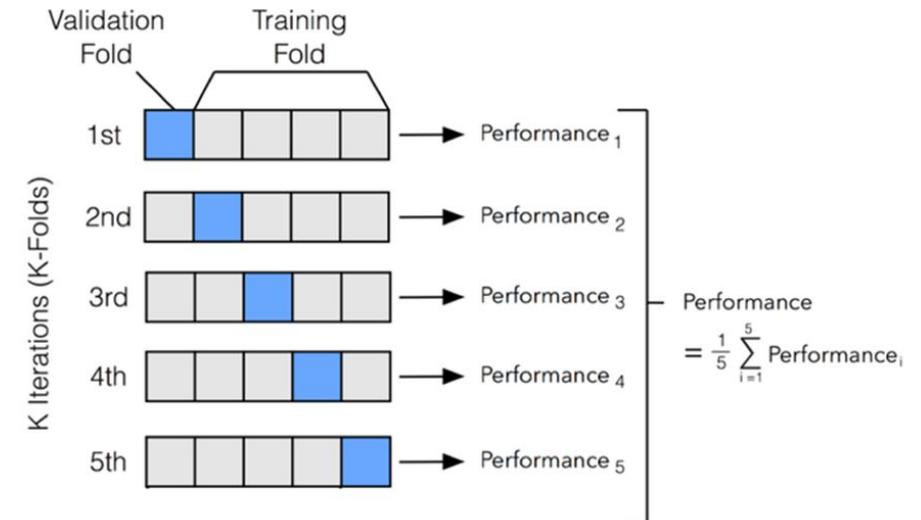


Figure 1-25. Model selection using holdout validation