



Introduction to Artificial Neural Networks with Keras

From Biological Neurons to Deep Learning with TensorFlow


From Biological to Artificial Neurons

- **Biological Neuron Inspiration**
 - Dendrites (inputs)
 - Cell body (processing)
 - Axon (output)
- **Artificial Neuron (Perceptron)**
 - Inputs: x_1, x_2, \dots, x_n
 - Weights: w_1, w_2, \dots, w_n
 - Bias: b
 - Activation Function: ϕ
 - Output: $y = \phi(w \cdot x + b)$



McCulloch Pitts Neuron

- $\hat{y} = \begin{cases} 1, & \text{if } \sum_{i=1}^n x_i \geq b \\ 0, & \text{otherwise} \end{cases}$
- $loss = \sum_i (y_i - \hat{y}_i)^2$
- Boolean input
- Boolean output
- Fixed slope
- Few possible intercepts (b)



The Perceptron

- Mathematical Model

$$z = w \cdot x + b$$

$$y = \phi(z)$$

- Step Function (Heaviside)

$$\phi(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{otherwise} \end{cases}$$


- Learning algorithm

- If $x \in P$ and $\sum_{i=1}^n w_i x_i < 0$

- $w = w + x$

- If $x \in N$ and $\sum_{i=1}^n w_i x_i \geq 0$

- $w = w - x$



Multilayer Perceptron (MLP)

- **Architecture:**


- Input Layer
- One or More Hidden Layers
- Output Layer

- **Feedforward Process**

$$a^{(l)} = \phi(W^{(l)}a^{(l-1)} + b^{(l)})$$

- **Universal Approximation**

Theorem: An MLP can approximate any continuous function



Backpropagation

- **Goal:** Minimize cost function $J(W, b)$
- **Process:**
 - Forward pass: compute output
 - Backward pass: compute gradients using chain rule
 - Update weights and biases using gradient descent
- **Gradient Update Rule:**

$$W \leftarrow W - \eta \frac{\partial J}{\partial W}$$

$$b \leftarrow b - \eta \frac{\partial J}{\partial b}$$

Chain Rule

- Let's take a simple **2-layer feedforward network**:

$$\begin{aligned}z_1 &= W_1 x + b_1 \\a_1 &= f(z_1) \\z_2 &= W_2 a_1 + b_2 \\\hat{y} &= g(z_2)\end{aligned}$$

- $\text{Loss}(L) = \text{Loss}(\hat{y}, y)$
- We need

$$\frac{\partial L}{\partial W_1}, \frac{\partial L}{\partial b_1}, \frac{\partial L}{\partial W_2}, \frac{\partial L}{\partial b_2}$$

Backpropagation Using Chain Rule

- - **Step 1: Output layer gradient**
 - $\delta_2 = \frac{\partial L}{\partial z_2} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2}$
 - $\frac{\partial L}{\partial W_2} = \delta_2 a_1^T$
 - $\frac{\partial L}{\partial b_2} = \delta_2$
 - **Step 2: Hidden layer gradient**
 - $\delta_1 = \frac{\partial L}{\partial z_1} = \frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial a_1} \cdot \frac{\partial a_1}{\partial z_1}$
 - $\frac{\partial L}{\partial W_1} = \delta_1 x^T$
 - $\frac{\partial L}{\partial b_1} = \delta_1$


Activation Functions

- Sigmoid
- Hyperbolic tangent function (Tanh)
- Rectified Linear Unit function (ReLU)

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$ReLU(z) = \max(0, z)$$



Regression and Classification MLPs

- Softmax (for output layer in classification):

$$e(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$



Hyperparameter	Binary classification	Multilabel binary classification	Multiclass classification
Input and hidden layers	Same as regression	Same as regression	Same as regression
# output neurons	1	1 per label	1 per class
Output layer activation	Logistic	Logistic	Softmax
Loss function	Cross entropy	Cross entropy	Cross entropy

Hyperparameter	Typical value
# input neurons	One per input feature (e.g., $28 \times 28 = 784$ for MNIST)
# hidden layers	Depends on the problem, but typically 1 to 5
# neurons per hidden layer	Depends on the problem, but typically 10 to 100
# output neurons	1 per prediction dimension
Hidden activation	ReLU (or SELU, see Chapter 11)
Output activation	None, or ReLU/softplus (if positive outputs) or logistic/tanh (if bounded outputs)
Loss function	MSE or MAE/Huber (if outliers)

Introduction to Keras

- <https://drive.google.com/file/d/17DTPHndXAczQS6cDLOffz0IcVJVwSmd3/view?usp=sharing>

Fine-Tuning Hyperparameters

- **Number of Hidden Layers:**
 - Start with 1–2, increase for complex problems
 - Better outcome by increasing the number of layers instead of the number of neurons per layer
- **Number of Neurons per Layer:**
 - Often set as a decreasing pyramid
- **Learning Rate:**
 - Use learning rate scheduling or adaptive optimizers (e.g., Adam)
- **Batch Size:**
 - Smaller batches → noisier updates, better generalization
- **Activation Functions:**
 - ReLU for hidden layers, Softmax for output in classification